

January 16, 2024

```
[1]: # Import other modules not related to PySpark
import os
import sys
import pandas as pd
from pandas import DataFrame
import numpy as np
import matplotlib.pyplot as plt
import matplotlib.ticker as mtick
import matplotlib
from mpl_toolkits.mplot3d import Axes3D
import math
from IPython.core.interactiveshell import InteractiveShell
from datetime import *
import statistics as stats
# This helps auto print out the items without explicitly using 'print'
InteractiveShell.ast_node_interactivity = "all"
%matplotlib inline
```

```
[2]: # Import PySpark related modules
import pyspark
from pyspark.rdd import RDD
from pyspark.sql import Row
from pyspark.sql import DataFrame
from pyspark.sql import SparkSession
from pyspark.sql import SQLContext
from pyspark.sql import functions
from pyspark.sql.functions import lit, desc, col, size, array_contains\
, isnan, udf, hour, array_min, array_max, countDistinct
from pyspark.sql.types import *

MAX_MEMORY = '15G'
# Initialize a spark session.
conf = pyspark.SparkConf().setMaster("local[*]") \
    .set('spark.executor.heartbeatInterval', 10000) \
    .set('spark.network.timeout', 10000) \
    .set("spark.core.connection.ack.wait.timeout", "3600") \
    .set("spark.executor.memory", MAX_MEMORY) \
```

```

        .set("spark.driver.memory", MAX_MEMORY)
def init_spark():
    spark = SparkSession \
        .builder \
        .appName("COVID") \
        .config(conf=conf) \
        .getOrCreate()
    return spark

spark = init_spark()
filename_data = 'the-reddit-covid-dataset-comments.csv'
# Load the main data set into pyspark data frame
df = spark.read.csv(filename_data, header=True, inferSchema=True)
print('Data frame type: ' + str(type(df)))

```

Data frame type: <class 'pyspark.sql.dataframe.DataFrame'>

```

[3]: #
df = df.sample(withReplacement=False, fraction=0.01, seed=42)
df = df.limit(2000)

```

```

[4]: print('Columns overview')
pd.DataFrame(df.dtypes, columns = ['Column Name', 'Data type'])

```

Columns overview

```

[4]:      Column Name Data type
0          type      string
1           id      string
2  subreddit.id      string
3  subreddit.name      string
4  subreddit.nsfw      string
5   created_utc      string
6    permalink      string
7         body      string
8    sentiment      string
9         score      string

```

```

[5]: #
new_column_names = ['type', 'id', 'subreddit_id', 'subreddit_name',
                    ↪ 'subreddit_nsfw', 'created_utc', 'permalink', 'body', 'sentiment', 'score']
#
df = df.toDF(*new_column_names)

pd.DataFrame(df.dtypes, columns = ['Column Name', 'Data type'])
df.printSchema()

```

```
[5]:      Column Name Data type
      0          type      string
      1           id      string
      2  subreddit_id      string
      3  subreddit_name      string
      4  subreddit_nsfw      string
      5   created_utc      string
      6   permalink      string
      7          body      string
      8   sentiment      string
      9          score      string
```

```
root
|-- type: string (nullable = true)
|-- id: string (nullable = true)
|-- subreddit_id: string (nullable = true)
|-- subreddit_name: string (nullable = true)
|-- subreddit_nsfw: string (nullable = true)
|-- created_utc: string (nullable = true)
|-- permalink: string (nullable = true)
|-- body: string (nullable = true)
|-- sentiment: string (nullable = true)
|-- score: string (nullable = true)
```

```
[6]: #
      selected_columns = ['id', 'subreddit_name', 'subreddit_nsfw', 'sentiment',
      ↪ 'score']
      df = df.select(selected_columns)

      #
      df.show(4)
```

```
+-----+-----+-----+-----+-----+
|      id|subreddit_name|subreddit_nsfw|sentiment|score|
+-----+-----+-----+-----+-----+
|  NULL|      NULL|      NULL|      NULL| NULL|
|hi1v4vl|      canada|      false| -0.7269|    1|
|hi1uyme| conservative|      false|      NULL| NULL|
|hi1udht|      jontron|      false|      NULL| NULL|
+-----+-----+-----+-----+-----+

only showing top 4 rows
```

```
[7]: df = df.na.drop(subset=['id', 'subreddit_name', 'subreddit_nsfw', 'sentiment',
      ↪ 'score'])
      df.show(2)
```

id	subreddit_name	subreddit_nsfw	sentiment	score
hi1v4vl	canada	false	-0.7269	1
hi1satq	nrl	false	-0.7506	46

only showing top 2 rows

```
[8]: # "sentiment"
df = df.withColumn("sentiment", col("sentiment").cast(FloatType()))

# "sentiment"
df = df.withColumn("score", col("score").cast(FloatType()))
```

```
[9]: #
quantiles = df.stat.approxQuantile(["sentiment", "score"], [0.25, 0.75], 0.05)

# (IQR)
IQR_sentiment = quantiles[0][1] - quantiles[0][0]
IQR_score = quantiles[1][1] - quantiles[1][0]

#
lower_bound_sentiment = quantiles[0][0] - 1.5 * IQR_sentiment
upper_bound_sentiment = quantiles[0][1] + 1.5 * IQR_sentiment

lower_bound_score = quantiles[1][0] - 1.5 * IQR_score
upper_bound_score = quantiles[1][1] + 1.5 * IQR_score

#
df_filtered = df.filter((col("sentiment").between(lower_bound_sentiment,
↪upper_bound_sentiment)) &
                        (col("score").between(lower_bound_score,
↪upper_bound_score)))

#
df_filtered.show(3)
```

id	subreddit_name	subreddit_nsfw	sentiment	score
hi1v4vl	canada	false	-0.7269	1.0
hi1q4qb	toiletpaperusa	false	0.4815	3.0
hi1pi1o	ukpolitics	false	-0.9432	1.0

only showing top 3 rows

```
[10]: #
summary = df_filtered.describe()

#
summary.show(3)
```

```
+-----+-----+-----+-----+-----+
---+
|summary| id|subreddit_name|subreddit_nsfw|          sentiment|
score|
+-----+-----+-----+-----+-----+
---+
| count| 246|          246|          246|          246|
246|
|  mean|NULL|          NULL|          NULL|          NULL|
NULL|-0.00225569022349...|2.0853658536585367|
| stddev|NULL|          NULL|          NULL|          NULL|
0.5579037134764742|1.9369735672689639|
+-----+-----+-----+-----+-----+
---+
only showing top 3 rows
```

```
[11]: from pyspark.sql import SparkSession
from pyspark.sql.functions import mean
#
mean_score = df_filtered.select(mean("score").alias("mean_score")).
    ↪collect()[0]["mean_score"]

#           (25%, 50%, 75%)           "score"
quantiles_score = df_filtered.stat.approxQuantile("score", [0.25, 0.5, 0.75], 0.
    ↪05)

#
print(f"Mean Score: {mean_score}")
print(f"25th Percentile: {quantiles_score[0]}")
print(f"50th Percentile (Median): {quantiles_score[1]}")
print(f"75th Percentile: {quantiles_score[2]}")
```

```
Mean Score: 2.0853658536585367
25th Percentile: 1.0
50th Percentile (Median): 1.0
75th Percentile: 2.0
```

```
[12]: #
mean_score = df_filtered.select(mean("sentiment").alias("mean_sentiment")).
    ↪collect()[0]["mean_sentiment"]
```

```

#               (25%, 50%, 75%)               "sentiment"
quantiles_score = df_filtered.stat.approxQuantile("sentiment", [0.25, 0.5, 0.
↳75], 0.05)

#
print(f"Mean Score: {mean_score}")
print(f"25th Percentile: {quantiles_score[0]}")
print(f"50th Percentile (Median): {quantiles_score[1]}")
print(f"75th Percentile: {quantiles_score[2]}")

```

Mean Score: -0.002255690223499527
 25th Percentile: -0.4018999934196472
 50th Percentile (Median): 0.0
 75th Percentile: 0.3400000035762787

```

[13]: from pyspark.sql.functions import col, count, lit, when, sum
#               subreddit
subreddit_counts = df_filtered.groupBy("subreddit_name").agg(count("*").
↳alias("count"))

#
total_count = df_filtered.count()

#
subreddit_percentages = subreddit_counts.withColumn(
    "percentage",
    (col("count") / total_count * 100).cast("double")
)

#               Other (      ,      ,      )
threshold = 1

#               Other
subreddit_percentages_filtered = (
    subreddit_percentages
    .withColumn("subreddit_grouped",
        when(col("count") > threshold, col("subreddit_name"))
        .otherwise(lit("Other")))
    .groupBy("subreddit_grouped")
    .agg(sum("count").alias("total_count"), sum("percentage").
↳alias("total_percentage"))
    .orderBy("total_percentage", ascending=False)
)

#               Pandas DataFrame
subreddit_percentages_filtered_pd = subreddit_percentages_filtered.toPandas()

```

```
#                                     truncate=False
print(subreddit_percentages_filtered_pd)
```

	subreddit_grouped	total_count	total_percentage
0	Other	128	52.032520
1	askreddit	12	4.878049
2	conspiracy	12	4.878049
3	hermancainaaward	12	4.878049
4	newzealand	5	2.032520
5	worldnews	5	2.032520
6	whitepeopletwitter	4	1.626016
7	antiwork	4	1.626016
8	politics	4	1.626016
9	soccer	4	1.626016
10	anarcho_capitalism	3	1.219512
11	libertarian	3	1.219512
12	byebyejob	3	1.219512
13	amitheasshole	3	1.219512
14	facepalm	3	1.219512
15	joerogan	3	1.219512
16	science	3	1.219512
17	louderwithcrowder	3	1.219512
18	ukpolitics	2	0.813008
19	movies	2	0.813008
20	unpopularopinion	2	0.813008
21	politicalhumor	2	0.813008
22	ireland	2	0.813008
23	debatevaccines	2	0.813008
24	lockdownskepticism	2	0.813008
25	tinder	2	0.813008
26	singapore	2	0.813008
27	sanjose	2	0.813008
28	stopdrinking	2	0.813008
29	politicalcompassmemes	2	0.813008
30	fantasyfootball	2	0.813008
31	askanamerican	2	0.813008
32	holup	2	0.813008
33	teenagers	2	0.813008

```
[21]: nsfw_column = "subreddit_nsfw"

#
df_filtered = df_filtered.withColumn(nsfw_column, col(nsfw_column).
↳ cast("boolean"))

#                                     NSFW
```

```

grouped_df = df_filtered.groupBy(nsfw_column).count()

#
nsfw_count = grouped_df.filter(col(nsfw_column) == True).select("count").
    ↪first()[0]
sfw_count = grouped_df.filter(col(nsfw_column) == False).select("count").
    ↪first()[0]

#
total_count = nsfw_count + sfw_count
nsfw_percentage = (nsfw_count / total_count) * 100
sfw_percentage = (sfw_count / total_count) * 100

#
print(f"NSFW: {nsfw_count}          , {nsfw_percentage:.2f}%")
print(f"SFW: {sfw_count}           , {sfw_percentage:.2f}%")

#
plt.show()

```

```

NSFW: 1          , 0.41%
SFW: 245         , 99.59%

```

```

[18]: #
grouped_df = df_filtered.groupBy(((col("sentiment") / 0.2).cast("int") * 0.2).
    ↪alias("sentiment_group")).count()

#
grouped_df = grouped_df.sort("sentiment_group")

#           Pandas DataFrame           matplotlib
pandas_df = grouped_df.toPandas()

#
plt.bar(pandas_df["sentiment_group"], pandas_df["count"], width=0.2)
plt.xlabel("Sentiment Group")
plt.ylabel("Count")
plt.title("Sentiment Analysis Distribution")
plt.show()


```

```
[18]: <BarContainer object of 9 artists>
```

```
[18]: Text(0.5, 0, 'Sentiment Group')
```

```
[18]: Text(0, 0.5, 'Count')
```

```
[18]: Text(0.5, 1.0, 'Sentiment Analysis Distribution')
```

```
[19]: #          subreddit_name      score      subreddit_name
subreddit_scores = df_filtered.groupBy("subreddit_name").sum("score")

#          score      -10
top_subreddits = subreddit_scores.sort(col("sum(score)").desc()).limit(10)

#
top_subreddits.show()
```

```
+-----+-----+
|  subreddit_name | sum(score) |
+-----+-----+
| hermancainaward |      46.0 |
|      conspiracy |      38.0 |
|      newzealand |      21.0 |
|      askreddit  |      18.0 |
| whitepeopletwitter |     13.0 |
|      worldnews  |      12.0 |
|      soccer     |      12.0 |
| lockdownskepticism |     10.0 |
```

```
|          ireland|      8.0|  
|lockdownskepticismaul|      8.0|  
+-----+-----+
```

```
[20]: #  
correlation = df_filtered.corr("sentiment", "score")  
  
#  
print(f"Correlation between sentiment and score: {correlation}")
```

Correlation between sentiment and score: -0.018082291461601126